Implementation with Ray (RLlib)

In last two weeks, I have

- implemented our environment with RLlib by modifying environments to fit RLlib standard input / output.
- conducted experiments using desktop computer using small batches.
- setup the cluster and configured Ray properly, with the help of Mingyang and cluster-managing stuffs.
- conducted experiments using cluster using large batches (learned how to use cluster for the first time…).

Performance of small batch training

Waiting and observing behaviors emerge! But still easy to collide.

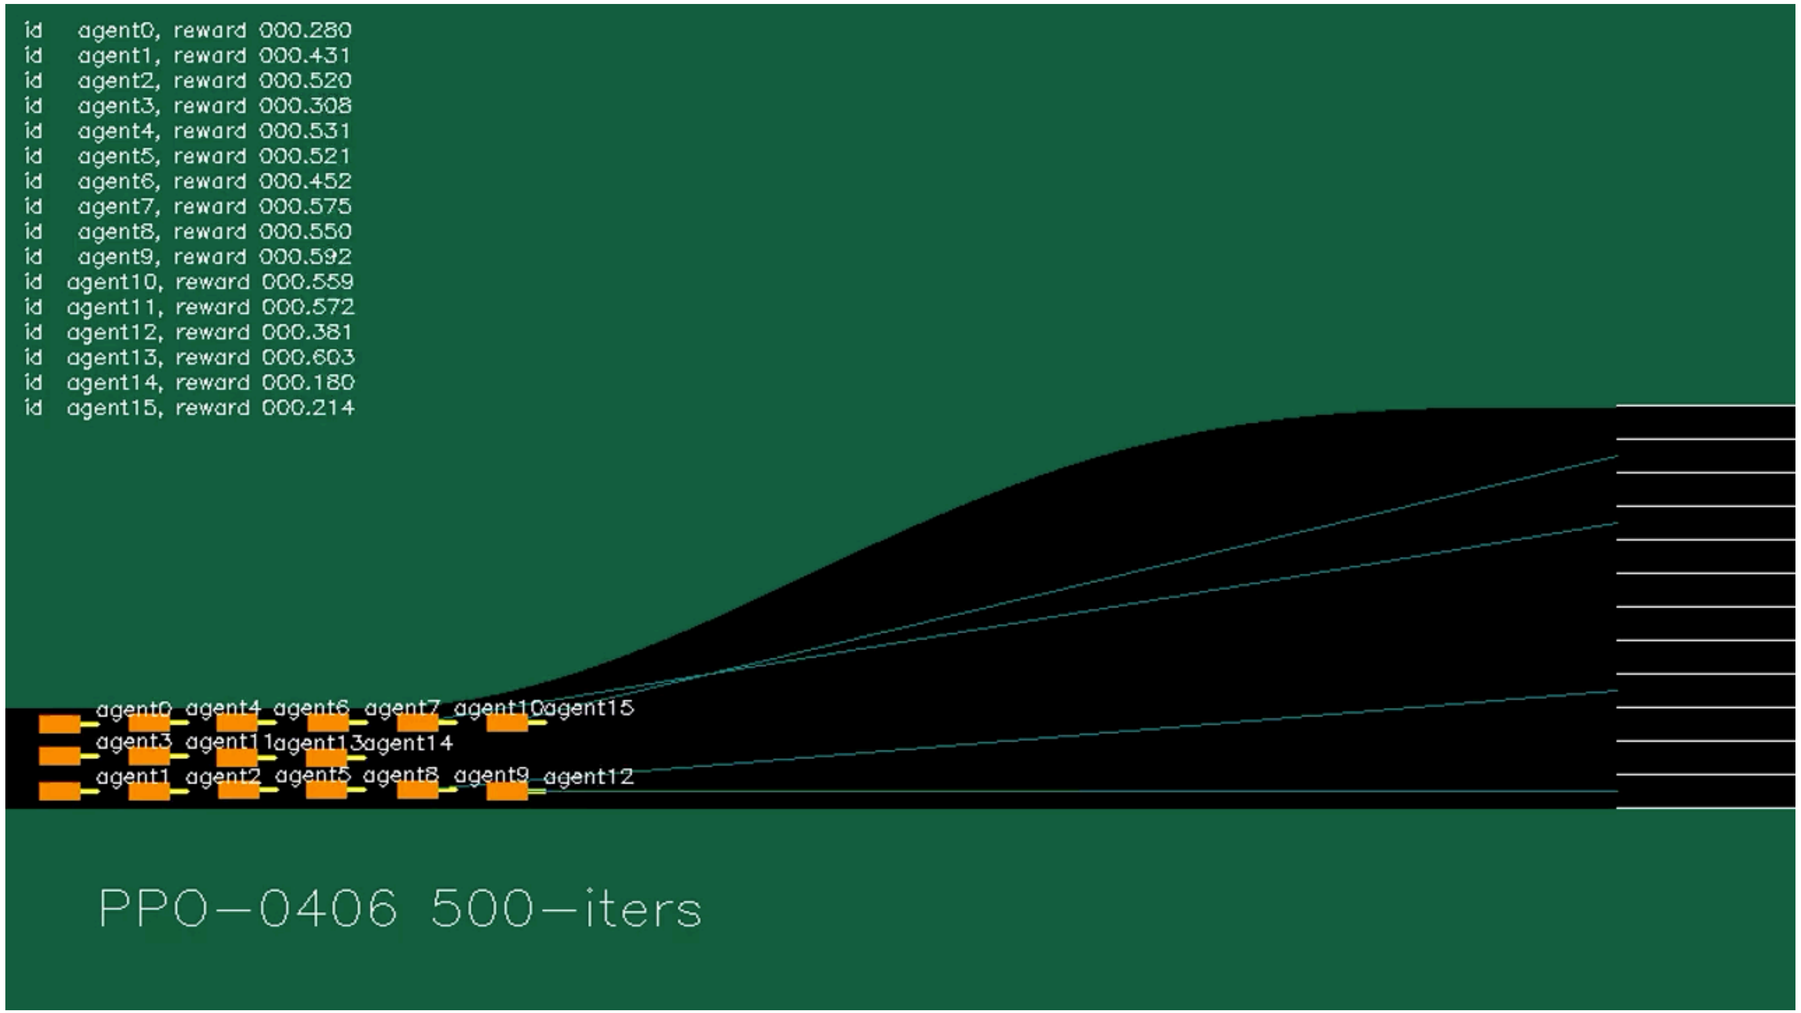

Metric	Value
iteration	500
— performance —	
agent collision	0.3075
boundary col. rate	0.0275
gate collision rate	0.1917
pass rate	0.4669
success rate	0.0756
episode reward	826.8
single reward	51.68
episode length	227.6
— configuration —	
num gpus	1
num workers	10
num env per worker	2
train batch size	4000